

```

import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error
from datetime import datetime
import math

```

```

train = pd.read_csv(r'C:\Users\kaler\Desktop\Untitled Folder 1\
train.csv')
features = pd.read_csv(r'C:\Users\kaler\Desktop\Untitled Folder 1\
features.csv')
stores = pd.read_csv(r'C:\Users\kaler\Desktop\Untitled Folder 1\
stores.csv')

```

```
train.head()
```

	Store	Dept	Date	Weekly_Sales	IsHoliday
0	1	1	2010-02-05	24924.50	False
1	1	1	2010-02-12	46039.49	True
2	1	1	2010-02-19	41595.55	False
3	1	1	2010-02-26	19403.54	False
4	1	1	2010-03-05	21827.90	False

```
stores.head()
```

	Store	Type	Size
0	1	A	151315
1	2	A	202307
2	3	B	37392
3	4	A	205863
4	5	B	34875

```
train.describe()
```

	Store	Dept	Weekly_Sales
count	421570.000000	421570.000000	421570.000000
mean	22.200546	44.260317	15981.258123
std	12.785297	30.492054	22711.183519
min	1.000000	1.000000	-4988.940000
25%	11.000000	18.000000	2079.650000
50%	22.000000	37.000000	7612.030000
75%	33.000000	74.000000	20205.852500
max	45.000000	99.000000	693099.360000

```
features.describe()
```

	Store	Temperature	Fuel_Price	MarkDown1
MarkDown2	\			

```

count    8190.000000    8190.000000    8190.000000    4032.000000
2921.000000
mean      23.000000     59.356198      3.405992      7032.371786
3384.176594
std       12.987966     18.678607      0.431337      9262.747448
8793.583016
min        1.000000     -7.290000      2.472000     -2781.450000    -
265.760000
25%       12.000000     45.902500      3.041000      1577.532500
68.880000
50%       23.000000     60.710000      3.513000      4743.580000
364.570000
75%       34.000000     73.880000      3.743000      8923.310000
2153.350000
max       45.000000    101.950000      4.468000    103184.980000
104519.540000

```

```

                Markdown3      Markdown4      Markdown5      CPI
Unemployment
count    3613.000000    3464.000000    4050.000000    7605.000000
7605.000000
mean     1760.100180    3292.935886    4132.216422    172.460809
7.826821
std      11276.462208    6792.329861    13086.690278    39.738346
1.877259
min      -179.260000      0.220000     -185.170000    126.064000
3.684000
25%       6.600000     304.687500    1440.827500    132.364839
6.634000
50%      36.260000    1176.425000    2727.135000    182.764003
7.806000
75%     163.150000    3310.007500    4832.555000    213.932412
8.567000
max    149483.310000    67474.850000    771448.100000    228.976456
14.313000

```

```
stores.describe()
```

```

                Store      Size
count    45.000000    45.000000
mean     23.000000    130287.600000
std      13.133926     63825.271991
min        1.000000    34875.000000
25%       12.000000    70713.000000
50%       23.000000    126512.000000
75%       34.000000    202307.000000
max       45.000000    219622.000000

```

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 421570 entries, 0 to 421569
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store           421570 non-null  int64
1   Dept            421570 non-null  int64
2   Date            421570 non-null  object
3   Weekly_Sales    421570 non-null  float64
4   IsHoliday       421570 non-null  bool
dtypes: bool(1), float64(1), int64(2), object(1)
memory usage: 13.3+ MB
```

```
features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8190 entries, 0 to 8189
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Store           8190 non-null   int64
1   Date            8190 non-null   object
2   Temperature     8190 non-null   float64
3   Fuel_Price      8190 non-null   float64
4   Markdown1       4032 non-null   float64
5   Markdown2       2921 non-null   float64
6   Markdown3       3613 non-null   float64
7   Markdown4       3464 non-null   float64
8   Markdown5       4050 non-null   float64
9   CPI             7605 non-null   float64
10  Unemployment     7605 non-null   float64
11  IsHoliday       8190 non-null   bool
dtypes: bool(1), float64(9), int64(1), object(1)
memory usage: 712.0+ KB
```

```
stores.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45 entries, 0 to 44
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Store   45 non-null     int64
1   Type    45 non-null     object
2   Size    45 non-null     int64
dtypes: int64(2), object(1)
memory usage: 1.2+ KB
```

```
train.isnull().sum()
```

```
Store      0
Dept       0
Date       0
Weekly_Sales  0
IsHoliday  0
dtype: int64
```

```
features.isnull().sum()
```

```
Store      0
Date       0
Temperature 0
Fuel_Price 0
MarkDown1  4158
MarkDown2  5269
MarkDown3  4577
MarkDown4  4726
MarkDown5  4140
CPI        585
Unemployment 585
IsHoliday  0
dtype: int64
```

```
stores.isnull().sum()
```

```
Store      0
Type       0
Size       0
dtype: int64
```

```
data = train.merge(features, on=['Store', 'Date'],
                    how='inner').merge(stores, on=['Store'],
                    how='inner')
print(data.shape)
```

```
(421570, 17)
```

```
data['MarkDown1'] = data['MarkDown1'].replace(np.nan,0)
data['MarkDown2'] = data['MarkDown2'].replace(np.nan,0)
data['MarkDown3'] = data['MarkDown3'].replace(np.nan,0)
data['MarkDown4'] = data['MarkDown4'].replace(np.nan,0)
data['MarkDown5'] = data['MarkDown5'].replace(np.nan,0)
```

```
data.describe()
```

	Store	Dept	Weekly_Sales	Temperature \
count	421570.000000	421570.000000	421570.000000	421570.000000
mean	22.200546	44.260317	15981.258123	60.090059
std	12.785297	30.492054	22711.183519	18.447931
min	1.000000	1.000000	-4988.940000	-2.060000
25%	11.000000	18.000000	2079.650000	46.680000

50%	22.000000	37.000000	7612.030000	62.090000
75%	33.000000	74.000000	20205.852500	74.280000
max	45.000000	99.000000	693099.360000	100.140000

	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	\
count	421570.000000	421570.000000	421570.000000	421570.000000	
mean	3.361027	2590.074819	879.974298	468.087665	
std	0.458515	6052.385934	5084.538801	5528.873453	
min	2.472000	0.000000	-265.760000	-29.100000	
25%	2.933000	0.000000	0.000000	0.000000	
50%	3.452000	0.000000	0.000000	0.000000	
75%	3.738000	2809.050000	2.200000	4.540000	
max	4.468000	88646.760000	104519.540000	141630.610000	

	MarkDown4	MarkDown5	CPI	Unemployment	\
count	421570.000000	421570.000000	421570.000000	421570.000000	
mean	1083.132268	1662.772385	171.201947	7.960289	
std	3894.529945	4207.629321	39.159276	1.863296	
min	0.000000	0.000000	126.064000	3.879000	
25%	0.000000	0.000000	132.022667	6.891000	
50%	0.000000	0.000000	182.318780	7.866000	
75%	425.290000	2168.040000	212.416993	8.572000	
max	67474.850000	108519.280000	227.232807	14.313000	

	Size
count	421570.000000
mean	136727.915739
std	60980.583328
min	34875.000000
25%	93638.000000
50%	140167.000000
75%	202505.000000
max	219622.000000

```
# Filter rows where 'weekly_Sales' is greater than or equal to 0
data_filtered = data[data['Weekly_Sales'] >= 0]
```

```
data.describe()
```

	Store	Dept	Weekly_Sales	Temperature	\
count	421570.000000	421570.000000	421570.000000	421570.000000	
mean	22.200546	44.260317	15981.258123	60.090059	
std	12.785297	30.492054	22711.183519	18.447931	
min	1.000000	1.000000	-4988.940000	-2.060000	
25%	11.000000	18.000000	2079.650000	46.680000	
50%	22.000000	37.000000	7612.030000	62.090000	
75%	33.000000	74.000000	20205.852500	74.280000	
max	45.000000	99.000000	693099.360000	100.140000	

	Fuel_Price	MarkDown1	MarkDown2	MarkDown3	\
count	421570.000000	421570.000000	421570.000000	421570.000000	
mean	3.361027	2590.074819	879.974298	468.087665	
std	0.458515	6052.385934	5084.538801	5528.873453	
min	2.472000	0.000000	-265.760000	-29.100000	
25%	2.933000	0.000000	0.000000	0.000000	
50%	3.452000	0.000000	0.000000	0.000000	
75%	3.738000	2809.050000	2.200000	4.540000	
max	4.468000	88646.760000	104519.540000	141630.610000	

	MarkDown4	MarkDown5	CPI	Unemployment	\
count	421570.000000	421570.000000	421570.000000	421570.000000	
mean	1083.132268	1662.772385	171.201947	7.960289	
std	3894.529945	4207.629321	39.159276	1.863296	
min	0.000000	0.000000	126.064000	3.879000	
25%	0.000000	0.000000	132.022667	6.891000	
50%	0.000000	0.000000	182.318780	7.866000	
75%	425.290000	2168.040000	212.416993	8.572000	
max	67474.850000	108519.280000	227.232807	14.313000	

	Size
count	421570.000000
mean	136727.915739
std	60980.583328
min	34875.000000
25%	93638.000000
50%	140167.000000
75%	202505.000000
max	219622.000000

```
data = pd.get_dummies(data,columns=['Type'])
```

```
data ['Date']= pd.to_datetime(data['Date'])
```

```
data['month'] = data['Date'].dt.month
```

```
data['Year'] = data['Date'].dt.year
```

```
data[['Date', 'month', 'Year']].head()
```

	Date	month	Year
0	2010-02-05	2	2010
1	2010-02-05	2	2010
2	2010-02-05	2	2010
3	2010-02-05	2	2010
4	2010-02-05	2	2010

```
data['dayofweek_name'] = data['Date'].dt.day_name()
```

```
data[['Date', 'dayofweek_name']].head()
```

	Date	dayofweek_name
0	2010-02-05	Friday

```

1 2010-02-05      Friday
2 2010-02-05      Friday
3 2010-02-05      Friday
4 2010-02-05      Friday

```

```

data['is_weekend'] = np.where(data['dayofweek_name'].isin(['Sunday',
'Saturday']), 1, 0)
data[['Date', 'is_weekend']].head()

```

```

      Date  is_weekend
0 2010-02-05         0
1 2010-02-05         0
2 2010-02-05         0
3 2010-02-05         0
4 2010-02-05         0

```

```

data["IsHoliday_x"] = data["IsHoliday_x"].astype(int)
del data['dayofweek_name']

print(data.head())

```

```

      Store Dept      Date  Weekly_Sales  IsHoliday_x  Temperature
Fuel_Price \
0      1      1 2010-02-05      24924.50          0          42.31
2.572
1      1      2 2010-02-05      50605.27          0          42.31
2.572
2      1      3 2010-02-05      13740.12          0          42.31
2.572
3      1      4 2010-02-05      39954.04          0          42.31
2.572
4      1      5 2010-02-05      32229.38          0          42.31
2.572

```

```

      Markdown1  Markdown2  Markdown3  ...      CPI  Unemployment \
0          0.0          0.0          0.0  ...  211.096358          8.106
1          0.0          0.0          0.0  ...  211.096358          8.106
2          0.0          0.0          0.0  ...  211.096358          8.106
3          0.0          0.0          0.0  ...  211.096358          8.106
4          0.0          0.0          0.0  ...  211.096358          8.106

```

```

      IsHoliday_y  Size  Type_A  Type_B  Type_C  month  Year
is_weekend
0      False  151315    True   False   False     2  2010
0
1      False  151315    True   False   False     2  2010
0
2      False  151315    True   False   False     2  2010
0
3      False  151315    True   False   False     2  2010

```

```

0
4      False  151315    True   False   False      2  2010
0

[5 rows x 22 columns]

data.to_csv('merged_data.csv',index=False)

df = pd.read_csv(r"C:\Users\kaler\Desktop\Untitled Folder 1\
merged_data.csv", keep_default_na=False, na_values=[""])
print(df.columns)

Index(['Store', 'Dept', 'Date', 'Weekly_Sales', 'IsHoliday_x',
      'Temperature',
      'Fuel_Price', 'MarkDown1', 'MarkDown2', 'MarkDown3',
      'MarkDown4',
      'MarkDown5', 'CPI', 'Unemployment', 'IsHoliday_y', 'Size',
      'Type_A',
      'Type_B', 'Type_C', 'month', 'Year', 'is_weekend'],
      dtype='object')

import pandas as pd
from sklearn.model_selection import train_test_split

# Assuming you have already read the CSV file into df

# Combine 'IsHoliday_x' and 'IsHoliday_y' into a single 'IsHoliday'
column
df['IsHoliday'] = df['IsHoliday_x'].combine_first(df['IsHoliday_y'])

# Check if 'IsHoliday' is in the DataFrame
if 'IsHoliday' in df.columns:
    # Select relevant columns in X
    X = df[["Store", "Dept", "Size", "IsHoliday", "CPI",
"Temperature", "Type_B", "Type_C", "MarkDown4", "month", "Year"]]

    # Create a copy of X to avoid the SettingWithCopyWarning
    X = X.copy()

    # Convert boolean values to float in 'Type_B' and 'Type_C'
    X['Type_B'] = X['Type_B'].astype(float)
    X['Type_C'] = X['Type_C'].astype(float)

    y = df['Weekly_Sales']

    # Reshape y to have a single column
    y = y.values.reshape(-1, 1)

    print(X.head())

```



```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
else:
    print("Column 'IsHoliday' not found in the DataFrame.")

```

Type_C \	Store	Dept	Size	IsHoliday	CPI	Temperature	Type_B
0	1	1	151315	0	211.096358	42.31	0.0
1	1	2	151315	0	211.096358	42.31	0.0
2	1	3	151315	0	211.096358	42.31	0.0
3	1	4	151315	0	211.096358	42.31	0.0
4	1	5	151315	0	211.096358	42.31	0.0

	Markdown4	month	Year
0	0.0	2	2010
1	0.0	2	2010
2	0.0	2	2010
3	0.0	2	2010
4	0.0	2	2010

```

pip install pmdarima

```

Defaulting to user installation because normal site-packages is not writeable

Requirement already satisfied: pmdarima in c:\users\kaler\appdata\roaming\python\python311\site-packages (2.0.4)

Requirement already satisfied: joblib>=0.11 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.2.0)

Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in c:\users\kaler\appdata\roaming\python\python311\site-packages (from pmdarima) (3.0.5)

Requirement already satisfied: numpy>=1.21.2 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.24.3)

Requirement already satisfied: pandas>=0.19 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (2.0.3)

Requirement already satisfied: scikit-learn>=0.22 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.3.0)

Requirement already satisfied: scipy>=1.3.2 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.11.1)

Requirement already satisfied: statsmodels>=0.13.2 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (0.14.0)

Requirement already satisfied: urllib3 in c:\programdata\anaconda3\lib\site-packages (from pmdarima) (1.26.16)

Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in c:\

```
programdata\anaconda3\lib\site-packages (from pmdarima) (68.0.0)
Requirement already satisfied: packaging>=17.1 in c:\programdata\
anaconda3\lib\site-packages (from pmdarima) (23.1)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\
programdata\anaconda3\lib\site-packages (from pandas>=0.19->pmdarima)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\
anaconda3\lib\site-packages (from pandas>=0.19->pmdarima)
(2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in c:\programdata\
anaconda3\lib\site-packages (from pandas>=0.19->pmdarima) (2023.3)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\
anaconda3\lib\site-packages (from scikit-learn>=0.22->pmdarima)
(2.2.0)
Requirement already satisfied: patsy>=0.5.2 in c:\programdata\
anaconda3\lib\site-packages (from statsmodels>=0.13.2->pmdarima)
(0.5.3)
Requirement already satisfied: six in c:\programdata\anaconda3\lib\
site-packages (from patsy>=0.5.2->statsmodels>=0.13.2->pmdarima)
(1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

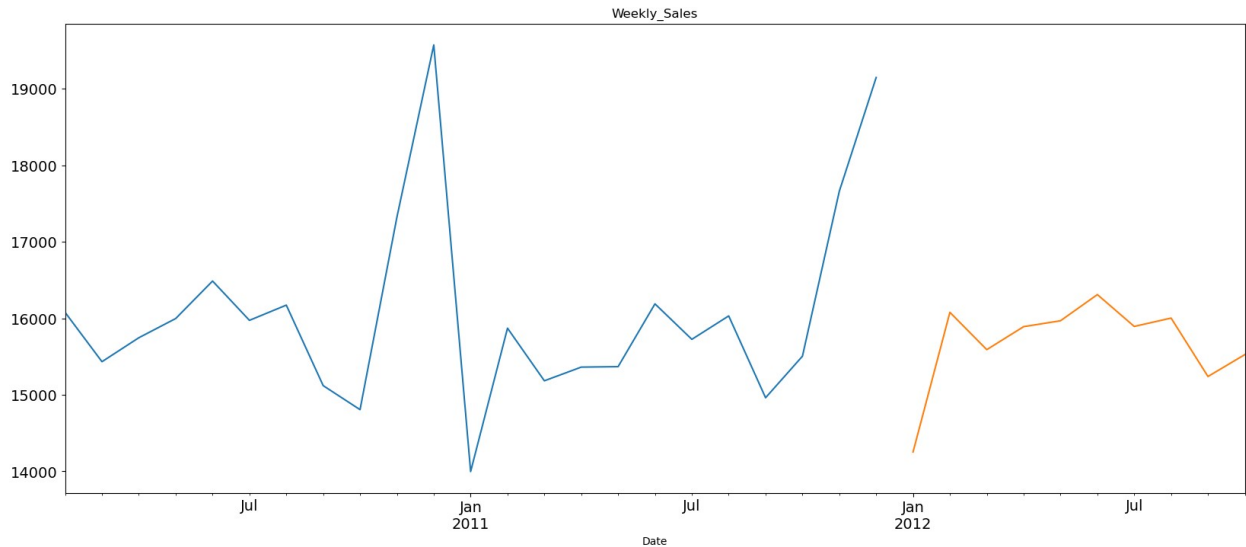
```
import pmdarima
from pmdarima.arima import auto_arima

df.Date = pd.to_datetime(df.Date, format='%Y-%m-%d')
df.index = df.Date
df = df.drop('Date', axis=1)
df = df.resample('MS').mean()

train_data = df[:int(0.7*(len(df)))]
test_data = df[int(0.7*(len(df))):]

train_data = train_data['Weekly_Sales']
test_data = test_data['Weekly_Sales']

train_data.plot(figsize=(20,8), title= 'Weekly_Sales', fontsize=14)
test_data.plot(figsize=(20,8), title='Weekly_Sales', fontsize =14)
plt.show()
```



```
from pmdarima import auto_arima

# Assuming you have a 'train_data' time series
model_auto_arima = auto_arima(train_data,
                               trace=True,
                               error_action='ignore',
                               suppress_warnings=True,
                               seasonal=True,
                               stepwise=False,
                               start_p=0, start_q=0,
                               max_p=10, max_q=10,
                               start_P=0, start_Q=0,
                               max_P=10, max_Q=10,
                               d=1, D=1, max_D=10,
                               approximation=False)

model_auto_arima.fit(train_data)
```

```
ARIMA(0,1,0)(0,0,0)[1] intercept : AIC=391.554, Time=0.03 sec
ARIMA(0,1,1)(0,0,0)[1] intercept : AIC=391.889, Time=0.04 sec
ARIMA(0,1,2)(0,0,0)[1] intercept : AIC=393.641, Time=0.04 sec
ARIMA(0,1,3)(0,0,0)[1] intercept : AIC=inf, Time=0.15 sec
ARIMA(0,1,4)(0,0,0)[1] intercept : AIC=397.671, Time=0.20 sec
ARIMA(0,1,5)(0,0,0)[1] intercept : AIC=inf, Time=0.38 sec
ARIMA(1,1,0)(0,0,0)[1] intercept : AIC=392.085, Time=0.05 sec
ARIMA(1,1,1)(0,0,0)[1] intercept : AIC=393.555, Time=0.15 sec
ARIMA(1,1,2)(0,0,0)[1] intercept : AIC=395.595, Time=0.13 sec
ARIMA(1,1,3)(0,0,0)[1] intercept : AIC=inf, Time=0.34 sec
ARIMA(1,1,4)(0,0,0)[1] intercept : AIC=inf, Time=0.38 sec
ARIMA(2,1,0)(0,0,0)[1] intercept : AIC=393.685, Time=0.04 sec
ARIMA(2,1,1)(0,0,0)[1] intercept : AIC=395.576, Time=0.07 sec
ARIMA(2,1,2)(0,0,0)[1] intercept : AIC=395.453, Time=0.39 sec
ARIMA(2,1,3)(0,0,0)[1] intercept : AIC=inf, Time=0.41 sec
```

```

ARIMA(3,1,0)(0,0,0)[1] intercept : AIC=395.636, Time=0.05 sec
ARIMA(3,1,1)(0,0,0)[1] intercept : AIC=397.572, Time=0.11 sec
ARIMA(3,1,2)(0,0,0)[1] intercept : AIC=inf, Time=0.34 sec
ARIMA(4,1,0)(0,0,0)[1] intercept : AIC=397.546, Time=0.06 sec
ARIMA(4,1,1)(0,0,0)[1] intercept : AIC=399.554, Time=0.12 sec
ARIMA(5,1,0)(0,0,0)[1] intercept : AIC=399.314, Time=0.07 sec

```

```

Best model: ARIMA(0,1,0)(0,0,0)[1] intercept
Total fit time: 3.587 seconds

```

```


ARIMA(order=(0, 1, 0), scoring_args={}, seasonal_order=(0, 0, 0, 1),
      suppress_warnings=True)

```

```

forecast = model_auto_arima.predict(n_periods=len(test_data))
forecast = pd.DataFrame(forecast, index =
test_data.index, columns=['Predictions'])
plt.figure(figsize=(20, 6))
plt.title('Prediction of Weekly Sales using Auto ARIMA model',
fontsize=20)
plt.plot(train_data, label='Train')
plt.plot(test_data, label='Test')
plt.plot(forecast, label='Prediction using ARIMA Model')
plt.legend(loc='best')
plt.xlabel('Date', fontsize=14)
plt.ylabel('Weekly Sales', fontsize=14)
plt.show()

```



```

import pandas as pd
import matplotlib.pyplot as plt
from pmdarima import auto_arima
from sklearn.metrics import mean_squared_error, mean_absolute_error
import numpy as np

# Assuming you have already defined 'train_data' and 'test_data'
# ... (Code for training the model and generating 'forecast')

```

```

# Use auto_arma for hyperparameter tuning
model_auto_arma = auto_arma(train_data, trace=True,
error_action='ignore', suppress_warnings=True, seasonal=True,
stepwise=False)

# Fit the auto_arma model
model_auto_arma.fit(train_data)

# Generate forecast
forecast = model_auto_arma.predict(n_periods=len(test_data))
forecast = pd.DataFrame(forecast, index=test_data.index,
columns=['Predictions'])

# Plotting
plt.figure(figsize=(20, 6))
plt.title('Prediction of Weekly Sales using Auto ARIMA model',
fontsize=20)
plt.plot(train_data, label='Train')
plt.plot(test_data, label='Test')
plt.plot(forecast, label='Prediction using ARIMA Model')
plt.legend(loc='best')
plt.xlabel('Date', fontsize=14)
plt.ylabel('Weekly Sales', fontsize=14)
plt.show()

# Calculate metrics
predictions = forecast['Predictions']
mse = mean_squared_error(test_data, predictions)
rmse = np.sqrt(mse)
mae = mean_absolute_error(test_data, predictions)

print('Mean Squared Error (MSE) of Auto ARIMA:', mse)
print('Root Mean Squared Error (RMSE) of Auto ARIMA:', rmse)
print('Mean Absolute Deviation (MAE) of Auto ARIMA:', mae)

ARIMA(0,0,0)(0,0,0)[1] intercept : AIC=397.842, Time=0.02 sec
ARIMA(0,0,1)(0,0,0)[1] intercept : AIC=399.420, Time=0.09 sec
ARIMA(0,0,2)(0,0,0)[1] intercept : AIC=inf, Time=0.18 sec
ARIMA(0,0,3)(0,0,0)[1] intercept : AIC=inf, Time=0.20 sec
ARIMA(0,0,4)(0,0,0)[1] intercept : AIC=399.645, Time=0.25 sec
ARIMA(0,0,5)(0,0,0)[1] intercept : AIC=inf, Time=0.41 sec
ARIMA(1,0,0)(0,0,0)[1] intercept : AIC=399.658, Time=0.03 sec
ARIMA(1,0,1)(0,0,0)[1] intercept : AIC=401.559, Time=0.06 sec
ARIMA(1,0,2)(0,0,0)[1] intercept : AIC=404.248, Time=0.09 sec
ARIMA(1,0,3)(0,0,0)[1] intercept : AIC=400.338, Time=0.13 sec
ARIMA(1,0,4)(0,0,0)[1] intercept : AIC=inf, Time=0.30 sec
ARIMA(2,0,0)(0,0,0)[1] intercept : AIC=399.687, Time=0.10 sec
ARIMA(2,0,1)(0,0,0)[1] intercept : AIC=403.287, Time=0.20 sec
ARIMA(2,0,2)(0,0,0)[1] intercept : AIC=404.649, Time=0.33 sec
ARIMA(2,0,3)(0,0,0)[1] intercept : AIC=402.435, Time=0.41 sec


```

```

ARIMA(3,0,0)(0,0,0)[1] intercept : AIC=400.869, Time=0.29 sec
ARIMA(3,0,1)(0,0,0)[1] intercept : AIC=403.022, Time=0.22 sec
ARIMA(3,0,2)(0,0,0)[1] intercept : AIC=405.264, Time=0.33 sec
ARIMA(4,0,0)(0,0,0)[1] intercept : AIC=402.820, Time=0.17 sec
ARIMA(4,0,1)(0,0,0)[1] intercept : AIC=405.025, Time=0.41 sec
ARIMA(5,0,0)(0,0,0)[1] intercept : AIC=404.311, Time=0.18 sec

```

Best model: ARIMA(0,0,0)(0,0,0)[1] intercept  
Total fit time: 4.405 seconds



Mean Squared Error (MSE) of Auto ARIMA: 468477.9539606969  
Root Mean Squared Error (RMSE) of Auto ARIMA: 684.4544937106461  
Mean Absolute Deviation (MAE) of Auto ARIMA: 446.8196095294599

```

from sklearn.ensemble import RandomForestRegressor
import pickle

# Assuming you have already defined and trained your
RandomForestRegressor 'rf'
rf = RandomForestRegressor(n_estimators=50, max_depth=20,
min_samples_split=3, min_samples_leaf=1)
# ... (train the model with your data)

# Save the trained model to a file
with open('final_model.pkl', 'wb') as model_file:
    pickle.dump(rf, model_file)

from sklearn.ensemble import RandomForestRegressor

rf = RandomForestRegressor(n_estimators=50, max_depth=20,
min_samples_split=3, min_samples_leaf=1)
rf.fit(X_train, y_train.ravel()) # Assuming X_train, y_train are your
training data
print('Accuracy:', rf.score(X_test, y_test.ravel()) * 100, '%')
y_pred = rf.predict(X_test)

```

```

from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error

rms = mean_squared_error(y_test, y_pred, squared=False)
print('RMSE:', rms)
print('MAE:', mean_absolute_error(y_test, y_pred))
print("Before fitting the model")
rf.fit(X_train, y_train.ravel()) # Assuming X_train, y_train are your
training data
print("After fitting the model")

# Add similar print statements at different stages

Accuracy: 96.4577508824995 %
RMSE: 4315.378610052086
MAE: 1648.949481735198
Before fitting the model
After fitting the model

print('Training Accuracy:', rf.score(X_train, y_train.ravel())*100, '%')

Training Accuracy: 99.127412543821 %

import xgboost as xgb
import warnings

import xgboost as xgb

xg_reg = xgb.XGBRegressor(objective='reg:squarederror', nthread=4,
                           n_estimators=500, max_depth=4,
                           learning_rate=0.5)
xg_reg.fit(X_train, y_train)

XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.5,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=4, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=500, n_jobs=None,
              nthread=4,
              num_parallel_tree=None, ...)

```

```

pred=xg_reg.predict(X_train)
y_pred=xg_reg.predict(X_test)

print('Accuracy:',xg_reg.score(X_test,y_test)*100,'%')

rms = mean_squared_error(y_test, y_pred, squared=False)
print('RMSE:',rms)

print('MAE:',mean_absolute_error(y_test, y_pred))

Accuracy: 93.20892619359017 %
RMSE: 5975.146796472817
MAE: 3035.7967659356923

print('Training Accuracy:',xg_reg.score(X_train,y_train)*100,'%')

Training Accuracy: 94.29593668521562 %

pip install prettytable

Defaulting to user installation because normal site-packages is not
writeableNote: you may need to restart the kernel to use updated
packages.

Requirement already satisfied: prettytable in c:\users\kaler\appdata\
roaming\python\python311\site-packages (3.9.0)
Requirement already satisfied: wcwidth in c:\programdata\anaconda3\
lib\site-packages (from prettytable) (0.2.5)

from prettytable import PrettyTable

tb = PrettyTable()
tb.field_names = ["Model","Training Accuracy","Testing
Accuracy","RMSE","MAE"]
tb.add_row(["Random Forest", 99.11, 96.77, 4055.83, 1651.13])
tb.add_row(["XgBoost", 94.23, 93.72, 5660.68, 3129.36])

print(tb)

+-----+-----+-----+-----+
+-----+
|      Model      | Training Accuracy | Testing Accuracy |  RMSE  |
MAE  |
+-----+-----+-----+-----+
+-----+
| Random Forest |      99.11      |      96.77      | 4055.83 |
1651.13 |
|      XgBoost  |      94.23      |      93.72      | 5660.68 |
3129.36 |
+-----+-----+-----+-----+
+-----+

```



```

from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

# Random Forest
rf = RandomForestRegressor(n_estimators=58, max_depth=27,
min_samples_split=3, min_samples_leaf=1)
rf.fit(X_train, y_train.ravel())
y_pred_rf = rf.predict(X_test)

# XGBoost
xg_reg = xgb.XGBRegressor(objective='reg:squarederror', nthread=4,
n_estimators=500, max_depth=4, learning_rate=0.5)
xg_reg.fit(X_train, y_train)
pred_xg = xg_reg.predict(X_train)
y_pred_xg = xg_reg.predict(X_test)

from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

# Random Forest
rf = RandomForestRegressor(n_estimators=58, max_depth=27,
min_samples_split=3, min_samples_leaf=1)
rf.fit(X_train, y_train.ravel())
y_pred_rf = rf.predict(X_test)

# XGBoost
xg_reg = xgb.XGBRegressor(objective='reg:squarederror', nthread=4,
n_estimators=500, max_depth=4, learning_rate=0.5)
xg_reg.fit(X_train, y_train)
pred_xg = xg_reg.predict(X_train)
y_pred_xg = xg_reg.predict(X_test)

cv=cross_val_score(xg_reg,X,y,cv=6)

np.mean(cv)

0.7387614591639351

pickle.dump(rf, open('newfinal_model.pkl','wb'))

```